



Malware detection classifier

Supervised Machine Learning Analysis

Document Overview: This document outlines the methodology, data analysis, model training, and evaluation results for classifying files as benign or malicious based on file entropy, permissions, and structural attributes.

Author: Nakedi Tumelo Peta

Date: January 15, 2026

Contents

1. Project Overview	3
2. Data Specification	3
2.1 Dataset Structure	3
2.2 Feature Dictionary	3
3. Exploratory Data Analysis (EDA)	4
3.1 Feature Distributions by Class.....	4
4. Methodology	6
4.1 Model Selection	6
4.2 Training Process	6
5. Model Evaluation & Results	7
5.1 Quantitative Metrics	7
5.2 Confusion Matrix Analysis	8
5.3 ROC Curve Analysis	8
6. Conclusion	9
Appendix A: Python Source Code	9

Malware Detection Classifier: Technical Documentation

1. Project Overview

This project implements a **Supervised Classification** algorithm (Decision Tree) to automatically detect potential malware. By analyzing file characteristics such as size, permissions, and entropy, the model classifies files into one of two categories:

Benign (Safe) or **Malware** (Malicious).

Objective: To develop a predictive tool that acts as a first line of defense by flagging suspicious files based on their metadata and structural attributes.

2. Data Specification

The model utilizes a dataset representing file artifacts, each labelled as either benign or malicious.

2.1 Dataset Structure

- Sample Size (N):** 200 observations (Split into 150 Training, 50 Testing).
- Class Balance:**
 - Benign (0): 66%
 - Malware (1): 34%

2.2 Feature Dictionary

Feature Name	Type	Range/Unit	Description
file_size	Float	KB (Log-normal)	Independent Variable. The physical size of the file on disk.
permissions	Integer	Categorical (0-2)	Independent Variable. Access rights: 0 (Read-only), 1 (Read-Write), 2 (Executable).
entropy	Float	1.0 - 8.0 bits/byte	Independent Variable. A measure of randomness. High entropy often indicates packed or encrypted code (common in malware).


num_imports	Integer	Count	Independent Variable. The number of external libraries or functions the file calls.
malware	Integer	Binary (0 or 1)	Target Variable. The label indicating file nature (1 = Malware, 0 = Benign).

3. Exploratory Data Analysis (EDA)


Visual analysis was conducted to understand feature distributions and class separation.

3.1 Feature Distributions by Class


- **File Size (Figure 1):** Malware samples (orange) tend to cluster in specific size ranges compared to benign files, which show a broader distribution.



- **Entropy (Figure 2):** There is a distinct spike in malware samples at high entropy levels (near 8.0). This confirms that high randomness (often due to obfuscation) is a strong indicator of malicious intent.



- **Imports (Figure 3):** Malware samples show different patterns in the number of imported functions compared to benign software, aiding in discrimination.




4. Methodology

4.1 Model Selection

- **Algorithm:** Decision Tree Classifier.
- **Hyperparameters:** max_depth=4 (Limited to prevent overfitting).
- **Reasoning:** Decision trees provide clear, interpretable "if-then" rules (e.g., "IF permissions=Executable AND entropy > 7 THEN Malware"), which is critical for understanding *why* a file was flagged.

4.2 Training Process

1. **Splitting:**
 - Training Set: 75% (150 samples)
 - Testing Set: 25% (50 samples)
2. **Fitting:** The model learned optimal split points for features like entropy and permissions to maximize the "Gini Impurity" reduction (separation of classes). See **Figure 4** for the visualized tree structure.



5. Model Evaluation & Results

The model was evaluated on the unseen Test Set (50 samples).


5.1 Quantitative Metrics

The model achieved high performance across all key metrics:

- **Accuracy: 98%** (Correctly classified 49 out of 50 files).
- **Precision: 1.0 (100%)** (When it predicted malware, it was *always* correct).
- **Recall: 94.1%** (It successfully caught 94% of the actual malware samples).
- **F1 Score: 0.97** (Harmonic mean of precision and recall).

5.2 Confusion Matrix Analysis

The confusion matrix (**Figure 5**) details the exact prediction counts:



	Predicted Benign	Predicted Malware
Actual Benign	33 (True Negatives)	0 (False Positives)
Actual Malware	1 (False Negative)	16 (True Positives)

- **Key Insight:** There were **zero False Positives**, meaning no safe files were accidentally flagged as malware. There was only **one False Negative**, meaning one malicious file slipped through.

5.3 ROC Curve Analysis

- **AUC Score:** **0.99** (Area Under the Curve).
- **Graph (Figure 6):** The blue line hugs the top-left corner, indicating near-perfect discrimination capability between classes.

6. Conclusion

The Decision Tree classifier proved highly effective for this dataset, particularly in leveraging **Entropy** and **Permissions** to identify malware. With an AUC of 0.99 and 100% precision, the model is suitable for deployment as a high-confidence filtering tool.

Next Steps

- Investigate the single False Negative case to understand which features caused the misclassification.
- Test the model against a larger, more diverse dataset to ensure robustness.\

Appendix A: Python Source Code

Description:

The following script was used to process the data, train the Decision Tree, and generate the evaluation metrics and plots.

Environment:

- Python 3.x
- Libraries: pandas, matplotlib, seaborn, sklearn

Code Listing:

```
#-----  
#Solving Malware classification Detection using Machine Learning  
#-----  
  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from sklearn.model_selection import train_test_split  
from sklearn.tree import DecisionTreeClassifier, plot_tree  
from sklearn.metrics import (  
    accuracy_score, precision_score, recall_score, f1_score,  
    confusion_matrix, roc_curve, auc  
)
```

```
#-----  
# Generate synthetic dataset  
#-----  
  
np.random.seed(42)  
n_samples = 200  
  
# Feature 1: file_size (in KB) - log-normal distribution to simulate varying sizes  
file_size = np.random.lognormal(mean=5, sigma=1, size=n_samples)  
  
# Feature 2: permissions (encoded as integer: 0=read-only, 1=read-write,  
2=executable)  
permissions = np.random.choice([0, 1, 2], size=n_samples, p=[0.4, 0.4, 0.2])  
  
#Feature 3: entropy - uniform between 1 and 8 bits per byte  
entropy = np.random.uniform(1, 8, size=n_samples)  
  
#Feature 4: num_imports - number of imported library calls (Poisson)  
num_imports = np.random.poisson(lam=20, size=n_samples)  
  
# Combine features into a DataFrame  
df_cls = pd.DataFrame({  
    'file_size': file_size,  
    'permissions': permissions,  
    'entropy': entropy,  
    'num_imports': num_imports  
})
```

```
# Target variable: malware (1) or benign (0)
#We'll assume higher entropy, execute, many imports -> more likely malware
```

```
malware_prob = (
    0.3 * (entropy / entropy.max()) +
    0.3 * (permissions == 2).astype(float) +
    0.3 * (num_imports / num_imports.max()) +
    0.1 * np.random.rand(n_samples)
)
```

```
y = (malware_prob > 0.5).astype(int)
df_cls['malware'] = y
```

```
#-----
# Inspect the dataset
#-----
```

```
print("First 5 rows of the dataset:")
print(df_cls.head(), "\n")
```

```
print("Feature descriptions:")
print("- file_size: Size of the file in KB (log-normal distribution).")
print("- permissions: 0=read-only, 1=read-write, 2=executable.")
print("- entropy: Measure of randomness in a file (1-8 bits per byte).")
print("- num_imports: Number of imported library calls.")
print("- malware: Target variable (1=malware, 0=benign).\n")
```

```
print("Class distribution:")
```

```
print(df_cls['malware'].value_counts(normalize=True), "\n")
```

```
#-----
```

```
# Visualize feature distributions by class
```

```
#-----
```

```
for feature in ["file_size", "entropy", "num_imports"]:
```

```
    plt.figure()
```

```
    plt.hist(
```

```
        [df_cls[df_cls["malware"]==0][feature], df_cls[df_cls["malware"]==1][feature]],
```

```
        bins=20,
```

```
        label=["Benign", "Malware"]
```

```
    )
```

```
    plt.xlabel(feature)
```

```
    plt.ylabel("Count")
```

```
    plt.title(f"Distribution of {feature} by Class")
```

```
    plt.legend()
```

```
    plt.show()
```

```
#-----
```

```
# Train-test split
```

```
#-----
```

```
x = df_cls[["file_size", "permissions", "entropy", "num_imports"]]
```

```
y = df_cls["malware"]
```

```
x_train, x_test, y_train, y_test = train_test_split(
```

```
    x, y, test_size=0.25, random_state=42, stratify=y
```

```
)
```

```
print ("Training samples:", x_train.shape[0])
print ("Testing samples:", x_test.shape[0], "\n")

#-----
# Fit Decision Tree Classifier
#-----

clf = DecisionTreeClassifier(max_depth=4, random_state=42)
clf.fit(x_train, y_train)

print ("Decision Tree trained with max_depth=4\n")

#-----
# Evaluate on test set
#-----

y_pred = clf.predict(x_test)
y_proba = clf.predict_proba(x_test)[:, 1]

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print ("Test Accuracy:", accuracy)
print ("Test Precision:", precision)
print ("Test Recall:", recall)
print ("Test F1 Score:", f1, "\n")
```

```
# Confusion Matrix

cm = confusion_matrix(y_test, y_pred)

print("Confusion Matrix: \n", cm, "\n")

#-----

# plot the tree structure

#-----

plt.figure(figsize=(12,6))

plot_tree(
    clf,
    feature_names=x.columns,
    class_names=["Benign", "Malware"],
    filled=False,
    rounded=True
)

plt.title("Decision Tree Structure")

plt.show()

#-----

# Confusion Matrix heatmap

#-----

plt.figure()

plt.imshow(cm, interpolation='nearest')

plt.title("Confusion Matrix Heatmap")

plt.colorbar()

tick_labels = ['Benign', 'Malware']

plt.xticks([0,1], tick_labels)
```

```
plt.yticks([0,1], tick_labels)
plt.xlabel('Predicted label')
plt.ylabel('True label')
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        plt.text(j, i, cm[i, j], ha = "center", va="center")
```

```
plt.show()
```

```
#-----
```

```
# ROC Curve
```

```
#-----
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_proba)
```

```
roc_auc = auc(fpr, tpr)
```

```
plt.figure()
```

```
plt.plot(fpr, tpr)
```

```
plt.plot([0, 1], [0, 1], linestyle='--')
```

```
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
```

```
plt.title(f'ROC Curve (AUC = {roc_auc:.2f})')
```

```
plt.show()
```